

Regression and Classification: A Comprehensive Analysis

Gabriele Durante

In this report, we aim to explore various statistical learning methodologies applied to regression and classification problems.

In the first part of the report, we examine linear regression models applied to aquatic toxicity data. This involves both traditional linear effects and alternative representations using dummy encoding, providing insight into how these transformations impact model fit and predictive accuracy. We apply techniques such as backward elimination and forward selection to perform variable selection, comparing criteria like AIC and BIC. Moreover, we introduce regularization techniques, specifically ridge regression, to address potential overfitting, followed by a discussion on parameter optimization via cross-validation and bootstrap methods. To extend our analysis to non-linear modeling, we apply generalized additive models (GAMs) with smoothing splines to capture more complex relationships within the data. Regression trees are explored, where we employ cost-complexity pruning to optimize tree size, emphasizing the balance between model complexity and prediction accuracy.

In the second part, we shift to a classification problem using the Pima Indians Diabetes dataset. Here, we experiment with k-nearest neighbors (k-NN), generalized additive models, classification trees, bagged trees, random forests, and neural networks. Each method is rigorously evaluated, using cross-validation techniques to assess generalization error and performance trade-offs. The report concludes with a comparative analysis of these models, offering insights into the advantages and limitations of different statistical learning approaches within the context of high-dimensional data.

Regression Model Comparison for Predicting Aquatic Toxicity

In this study, we investigate the acute aquatic toxicity of various organic molecules, as quantified by the lethal concentration (LC50) that induces mortality in 50% of the planktonic crustacean *Daphnia magna* over a 48-hour exposure period. The objective is to develop predictive models that leverage molecular descriptors to ascertain the toxicity levels of these compounds. To this end, we utilize a dataset comprising 546 observations, which is sourced from the UCI Machine Learning Repository. The dataset encompasses eight key molecular descriptors identified as significant predictors of LC50:

- TPSA: Topological Polar Surface Area, calculated via a contribution method considering the presence of nitrogen, oxygen, potassium, and sulfur.
- SAacc: Van der Waals Surface Area (VSA) of hydrogen bond acceptor atoms.
- H050: The count of hydrogen atoms bonded to heteroatoms.
- MLOGP: A measure of lipophilicity, serving as a critical determinant of narcosis.
- RDCHI: A topological index encapsulating information related to molecular size and branching.
- GATS1p: An indicator of molecular polarizability.
- nN: The total number of nitrogen atoms within the molecule.
- C040: The count of specific carbon atom types, including those found in esters, carboxylic acids, thioesters, carbamic acids, and nitriles.

```
summary(data)
```

##	TPSA	SAacc	H050	MLOGP
##	Min. : 0.00	Min. : 0.00	Min. : 0.0000	Min. : -6.446
##	1st Qu.: 15.79	1st Qu.: 11.00	1st Qu.: 0.0000	1st Qu.: 1.232

```
## Median : 40.46      Median : 42.68      Median : 0.0000      Median : 2.273
## Mean   : 48.47      Mean   : 58.87      Mean   : 0.9377      Mean   : 2.313
## 3rd Qu.: 70.02      3rd Qu.: 77.49      3rd Qu.: 1.0000      3rd Qu.: 3.393
## Max.   :347.32      Max.   :571.95      Max.   :18.0000      Max.   : 9.148
##      RDCHI          GATS1p          nN          C040
## Min.   :1.000      Min.   :0.281      Min.   : 0.000      Min.   : 0.0000
## 1st Qu.:1.975      1st Qu.:0.737      1st Qu.: 0.000      1st Qu.: 0.0000
## Median :2.344      Median :1.020      Median : 1.000      Median : 0.0000
## Mean   :2.492      Mean   :1.046      Mean   : 1.004      Mean   : 0.3535
## 3rd Qu.:2.911      3rd Qu.:1.266      3rd Qu.: 2.000      3rd Qu.: 0.0000
## Max.   :6.439      Max.   :2.500      Max.   :11.000      Max.   :11.0000
##      LC50
## Min.   : 0.122
## 1st Qu.: 3.602
## Median : 4.516
## Mean   : 4.658
## 3rd Qu.: 5.607
## Max.   :10.047
```

```
colSums(is.na(data))
```

```
##      TPSA      SAacc      H050      MLOGP      RDCHI      GATS1p      nN      C040      LC50
##       0         0         0         0         0         0         0         0         0
```

With a complete dataset comprising 546 observations, we can ensure that our modeling efforts are not hindered by data gaps. This completeness facilitates a more accurate assessment of the relationships between the molecular descriptors and LC50, allowing for a more confident interpretation of the findings and their implications in predicting aquatic toxicity.

Linear and Dummy Encoding Approach

To evaluate the predictive models effectively, we partition the dataset into training and testing subsets. A random seed is set using `set.seed(2024)` to ensure reproducibility. The training set is constructed by sampling approximately two-thirds of the total observations with the command `index <- sample(1:nrow(data), size = round(2/3 * nrow(data)))`.

- Training Data: `trainData <- data[index,]` comprises the selected observations for model fitting.
- Testing Data: `testData <- data[-index,]` contains the remaining observations for model evaluation.

```
# build the dataset for the training
set.seed(2024)
index <- sample(1:nrow(data), size = round(2/3 * nrow(data)))
trainData <- data[index, ]
testData <- data[-index, ]
```

Linear Model

First we fitted a linear regression model to predict LC50 using molecular descriptors from the training dataset. The model's performance is summarized as follows:

```
set.seed(2024)
### Modeling Count Variables Directly as Linear Effects
modell1 <- lm(LC50 ~., data = trainData)

pred_model1_train <- predict(modell1, newdata = trainData)
pred_model1_test <- predict(modell1, newdata = testData)
```

```
error_train <- mean((pred_model1_train - trainData$LC50)^2)
error_test <- mean((pred_model1_test - testData$LC50)^2)
```

```
cat("Training Error (Linear):", error_train, "\n")
```

```
## Training Error (Linear): 1.377162
```

```
cat("Test Error (Linear):", error_test, "\n")
```

```
## Test Error (Linear): 1.556978
```

```
summary(model1)
```

```
##
## Call:
## lm(formula = LC50 ~ ., data = trainData)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4.1890 -0.7482 -0.1153  0.6079  3.7987
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  2.875587   0.298304   9.640 < 2e-16 ***
## TPSA         0.026589   0.003220   8.258 2.96e-15 ***
## SAacc        -0.012127   0.002522  -4.809 2.25e-06 ***
## H050          0.031343   0.070945   0.442  0.65891
## MLOGP         0.496920   0.077134   6.442 3.83e-10 ***
## RDCHI         0.308016   0.166170   1.854  0.06462 .
## GATS1p        -0.537871   0.187291  -2.872  0.00433 **
## nN            -0.198318   0.057041  -3.477  0.00057 ***
## C040          -0.055902   0.090200  -0.620  0.53582
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.188 on 355 degrees of freedom
## Multiple R-squared:  0.4721, Adjusted R-squared:  0.4602
## F-statistic: 39.68 on 8 and 355 DF, p-value: < 2.2e-16
```

The R-squared value of 0.472 indicates that while the linear model captures some of the variability in the data, it leaves a considerable amount unexplained. Highly significant predictors, such as TPSA ($p < 2e-16$) and MLOGP ($p < 3.83e-10$), reveal strong relationships with LC50, indicating that increased topological polar surface area and greater lipophilicity correlate with higher toxicity levels.

Dummy Linear Model

After, model dummy encoding was applied to the count variables **nN**, **C040**, and **H050** in both the training and test datasets to transform these continuous variables into binary indicators. Specifically, values greater than zero were recoded to 1, while values of zero were recoded to 0. This approach allows the linear model to assess the presence or absence of these variables as factors influencing acute aquatic toxicity. Predictions were generated for both the training and test datasets, enabling the calculation of the training error and test error.

```
### Dummy Encoding for Count Variables
```

```
train_dummy <- trainData
```

```
test_dummy <- testData
```

```
# dummy encoding
```

```

train_dummy$nN <- ifelse(train_dummy$nN > 0, 1, 0)
test_dummy$nN <- ifelse(test_dummy$nN > 0, 1, 0)
test_dummy$C040 <- ifelse(test_dummy$C040 > 0, 1, 0)
train_dummy$C040 <- ifelse(train_dummy$C040 > 0, 1, 0)
test_dummy$H050 <- ifelse(test_dummy$H050 > 0, 1, 0)
train_dummy$H050 <- ifelse(train_dummy$H050 > 0, 1, 0)

model2 <- lm(train_dummy$LC50 ~ ., data = train_dummy)

pred_model2_train <- predict(model2, newdata = train_dummy)
pred_model2_test <- predict(model2, newdata = test_dummy)

error_train_dummy <- mean((pred_model2_train - train_dummy$LC50)^2)
error_test_dummy <- mean((pred_model2_test - test_dummy$LC50)^2)

cat("Training Error (Dummy):", error_train_dummy, "\n")

## Training Error (Dummy): 1.424236

cat("Test Error (Dummy):", error_test_dummy, "\n")

## Test Error (Dummy): 1.616825

summary(model2)

##
## Call:
## lm(formula = train_dummy$LC50 ~ ., data = train_dummy)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4.0143 -0.7807 -0.1313  0.6313  3.7856
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  3.007963   0.305982   9.831 < 2e-16 ***
## TPSA         0.022380   0.003185   7.026 1.09e-11 ***
## SAacc        -0.010184   0.002193  -4.643 4.83e-06 ***
## H050         -0.097746   0.154467  -0.633 0.52727
## MLOGP        0.502377   0.076886   6.534 2.22e-10 ***
## RDCHI        0.262639   0.167774   1.565 0.11837
## GATS1p       -0.559217   0.185302  -3.018 0.00273 **
## nN           -0.054983   0.149020  -0.369 0.71237
## C040         -0.162055   0.162002  -1.000 0.31783
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.208 on 355 degrees of freedom
## Multiple R-squared:  0.454, Adjusted R-squared:  0.4417
## F-statistic: 36.9 on 8 and 355 DF, p-value: < 2.2e-16

```

The training error for the dummy-encoded model is 1.424, while the test error is 1.617. Compared to the previous linear model, which yielded a training error of 1.377 and a test error of 1.557, this model exhibits slightly higher error rates on both the training and test datasets. This increase suggests that the dummy encoding may not have improved the model's predictive performance. This could indicate a potential loss of information due to the transformation of continuous variables into binary indicators, which could limit the

model's ability to capture the nuances of the underlying relationships between predictors and the response variable.

The overall model performance is reflected in the Multiple R-squared value of 0.454, while the adjusted R-squared of 0.4417 accounts for the number of predictors in the model. The F-statistic (36.9) with a p-value $< 2.2\text{e-}16$ indicates that the model is statistically significant.

Empirical Error Distribution Analysis in Regression Models

The function `perform_analysis` is defined to conduct repeated evaluations of the two linear models defined before on the same dataset and returns the values of error test and train. The analysis is repeated 200 times using the `replicate` function, enabling the computation of average test errors for both models. The results are then visualized using a density plot, illustrating the empirical distribution of test errors for each model.

```
library(ggplot2)

perform_analysis <- function(data) {
  index <- sample(1:nrow(data), size = round(2/3 * nrow(data)))
  trainData <- data[index, ]
  testData <- data[-index, ]

  # model 1
  model1 <- lm(LC50 ~., data = trainData)
  pred_model1_train_200 <- predict(model1, newdata = trainData)
  pred_model1_test_200 <- predict(model1, newdata = testData)
  error_train_200_lm <- mean((pred_model1_train_200 - trainData$LC50)^2)
  error_test_200_lm <- mean((pred_model1_test_200 - testData$LC50)^2)

  # model 2
  train_dummy <- trainData
  test_dummy <- testData

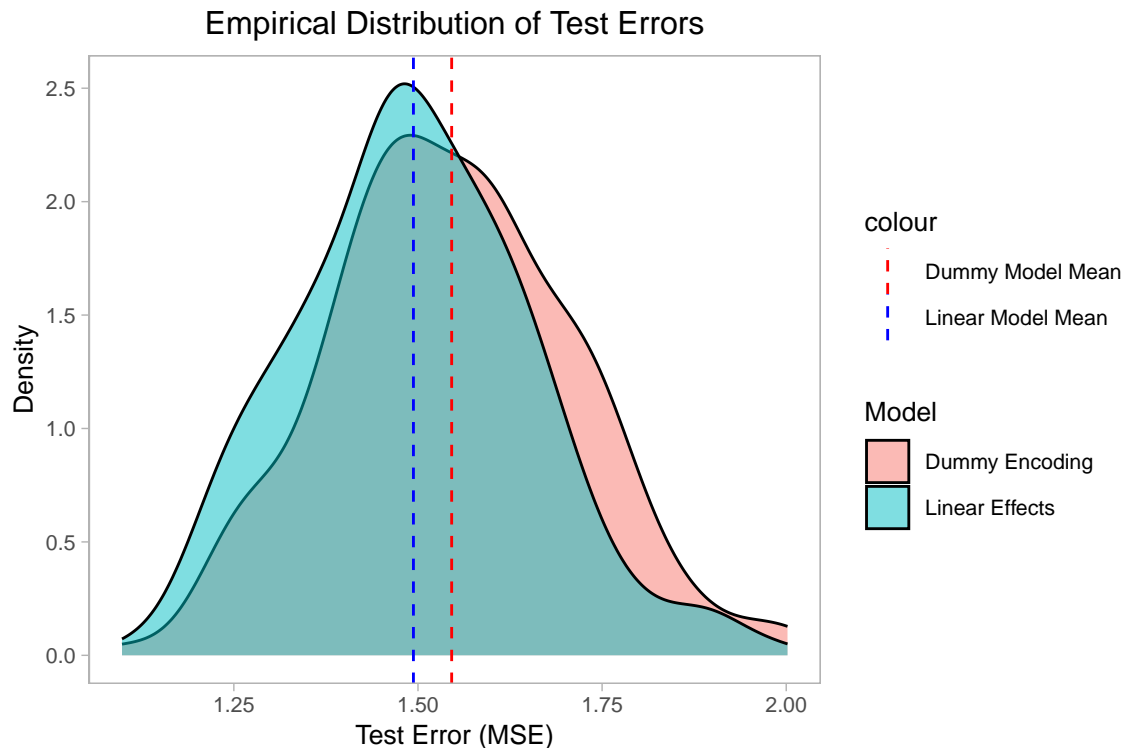
  # dummy encoding
  train_dummy$nN <- ifelse(train_dummy$nN > 0, 1, 0)
  test_dummy$nN <- ifelse(test_dummy$nN > 0, 1, 0)
  test_dummy$C040 <- ifelse(test_dummy$C040 > 0, 1, 0)
  train_dummy$C040 <- ifelse(train_dummy$C040 > 0, 1, 0)
  test_dummy$H050 <- ifelse(test_dummy$H050 > 0, 1, 0)
  train_dummy$H050 <- ifelse(train_dummy$H050 > 0, 1, 0)

  model2 <- lm(train_dummy$LC50 ~ ., data = train_dummy)

  pred_model2_train <- predict(model2, newdata = train_dummy)
  pred_model2_test <- predict(model2, newdata = test_dummy)

  error_train_200_dummy <- mean((pred_model2_train - train_dummy$LC50)^2)
  error_test_200_dummy <- mean((pred_model2_test - test_dummy$LC50)^2)

  return(c(error_test_200_lm, error_test_200_dummy))
}
```



The graph presents a density plot comparing the test error distributions for the two models. The aim should be illustrate the performance of the models on unseen data and their reliability, focusing on the variability and central tendency of their respective errors. The Linear Model exhibits a narrower distribution centered around a lower MSE compared to the Dummy Model, whose distribution appears wider and shifted slightly to the right, indicating higher average test error. Basically, the tighter distribution and lower mean MSE of the Linear Model suggest that it generalize better to the test set compared to the dummy model.

Variable Selection Methods in Linear Regression

```
library(MASS)
library(leaps)

# Backward Elimination
full_model <- lm(LC50 ~ ., data = trainData)
backward_aic <- stepAIC(full_model, direction = "backward",
                        k = 2, trace = FALSE)
summary(backward_aic)
```

```
##
## Call:
## lm(formula = LC50 ~ TPSA + SAacc + MLOGP + RDCHI + GATS1p + nN,
##     data = trainData)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4.1632 -0.7485 -0.1143  0.6156  3.8169
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
```

```
## (Intercept)  2.970390    0.278237   10.676 < 2e-16 ***
## TPSA         0.026358    0.003125    8.435 8.35e-16 ***
## SAacc        -0.011638    0.001946   -5.980 5.40e-09 ***
## MLOGP         0.487552    0.073711    6.614 1.37e-10 ***
## RDCHI         0.290699    0.162996    1.783 0.075358 .
## GATS1p        -0.573581    0.179110   -3.202 0.001485 **
## nN           -0.197504    0.054796   -3.604 0.000357 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.187 on 357 degrees of freedom
## Multiple R-squared:  0.4707, Adjusted R-squared:  0.4618
## F-statistic: 52.91 on 6 and 357 DF,  p-value: < 2.2e-16

backward_bic <- stepAIC(full_model, direction = "backward",
                        k = log(nrow(trainData)), trace = FALSE)
summary(backward_bic)

##
## Call:
## lm(formula = LC50 ~ TPSA + SAacc + MLOGP + GATS1p + nN, data = trainData)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4.0473 -0.7751 -0.0928  0.5879  3.7865
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  3.123831    0.265406   11.770 < 2e-16 ***
## TPSA         0.028625    0.002863    9.997 < 2e-16 ***
## SAacc        -0.010068    0.001741   -5.783 1.6e-08 ***
## MLOGP         0.594639    0.042886   13.866 < 2e-16 ***
## GATS1p       -0.466421    0.169244   -2.756 0.006152 **
## nN           -0.187790    0.054690   -3.434 0.000665 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.19 on 358 degrees of freedom
## Multiple R-squared:  0.466, Adjusted R-squared:  0.4585
## F-statistic: 62.47 on 5 and 358 DF,  p-value: < 2.2e-16

# Forward Selection
null_model <- lm(LC50 ~ 1, data = trainData)
forward_aic <- stepAIC(null_model, direction = "forward",
                      scope = formula(full_model), k = 2, trace = FALSE)
summary(forward_aic)

##
## Call:
## lm(formula = LC50 ~ MLOGP + TPSA + SAacc + nN + GATS1p + RDCHI,
##     data = trainData)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4.1632 -0.7485 -0.1143  0.6156  3.8169
```

```
##
## Coefficients:
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept)  2.970390   0.278237  10.676 < 2e-16 ***
## MLOGP        0.487552   0.073711   6.614 1.37e-10 ***
## TPSA         0.026358   0.003125   8.435 8.35e-16 ***
## SAacc        -0.011638   0.001946  -5.980 5.40e-09 ***
## nN           -0.197504   0.054796  -3.604 0.000357 ***
## GATS1p       -0.573581   0.179110  -3.202 0.001485 **
## RDCHI        0.290699   0.162996   1.783 0.075358 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.187 on 357 degrees of freedom
## Multiple R-squared:  0.4707, Adjusted R-squared:  0.4618
## F-statistic: 52.91 on 6 and 357 DF,  p-value: < 2.2e-16

forward_bic <- stepAIC(null_model, direction = "forward",
                      scope = formula(full_model), k = log(nrow(trainData)), trace = FALSE)
summary(forward_bic)

##
## Call:
## lm(formula = LC50 ~ MLOGP + TPSA + SAacc + nN + GATS1p, data = trainData)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4.0473 -0.7751 -0.0928  0.5879  3.7865
##
## Coefficients:
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept)  3.123831   0.265406  11.770 < 2e-16 ***
## MLOGP        0.594639   0.042886  13.866 < 2e-16 ***
## TPSA         0.028625   0.002863   9.997 < 2e-16 ***
## SAacc        -0.010068   0.001741  -5.783 1.6e-08 ***
## nN           -0.187790   0.054690  -3.434 0.000665 ***
## GATS1p       -0.466421   0.169244  -2.756 0.006152 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.19 on 358 degrees of freedom
## Multiple R-squared:  0.466, Adjusted R-squared:  0.4585
## F-statistic: 62.47 on 5 and 358 DF,  p-value: < 2.2e-16

# Compare Models
cat("Backward Elimination (AIC) Model:\n")

## Backward Elimination (AIC) Model:
print(backward_aic$call)

## lm(formula = LC50 ~ TPSA + SAacc + MLOGP + RDCHI + GATS1p + nN,
##      data = trainData)
cat("\nBackward Elimination (BIC) Model:\n")

##
```



```

## Backward Elimination (BIC) Model:
print(backward_bic$call)

## lm(formula = LC50 ~ TPSA + SAacc + MLOGP + GATS1p + nN, data = trainData)
cat("\nForward Selection (AIC) Model:\n")

##
## Forward Selection (AIC) Model:
print(forward_aic$call)

## lm(formula = LC50 ~ MLOGP + TPSA + SAacc + nN + GATS1p + RDCHI,
##     data = trainData)
cat("\nForward Selection (BIC) Model:\n")

##
## Forward Selection (BIC) Model:
print(forward_bic$call)

## lm(formula = LC50 ~ MLOGP + TPSA + SAacc + nN + GATS1p, data = trainData)
# Model Comparisons
cat("\nModels Summary Comparison:\n")

##
## Models Summary Comparison:
cat("AIC Backward:", AIC(backward_aic), "\n")

## AIC Backward: 1166.421
cat("AIC Forward:", AIC(forward_aic), "\n")

## AIC Forward: 1166.421
cat("BIC Backward:", BIC(backward_bic), "\n")

## BIC Backward: 1194.93
cat("BIC Forward:", BIC(forward_bic), "\n")

## BIC Forward: 1194.93

```

Ridge Regression Model Optimization

```

library(glmnet)
library(boot)

index <- sample(seq_len(nrow(data)), size = 2/3 * nrow(data))
trainData <- data[index, ]
testData <- data[-index, ]

X_train <- as.matrix(trainData[, -ncol(trainData)])
y_train <- trainData$LC50
X_test <- as.matrix(testData[, -ncol(testData)])
y_test <- testData$LC50

```

```

lambda_grid <- 10^seq(3, -2, length = 100)

# Ridge regression (CV)
cv_ridge <- cv.glmnet(X_train, y_train, alpha = 0,
                      lambda = lambda_grid, nfolds = 10)
optimal_lambda_cv <- cv_ridge$lambda.min

cat("Optimal lambda from Cross-Validation:", optimal_lambda_cv, "\n")

## Optimal lambda from Cross-Validation: 0.01

ridge_model <- glmnet(X_train, y_train, alpha = 0,
                      lambda = optimal_lambda_cv)
train_predictions_ridge <- predict(ridge_model, newx = X_train)
test_predictions_ridge <- predict(ridge_model, newx = X_test)
mse_train_ridge <- mean((y_train - train_predictions_ridge)^2)
mse_test_ridge <- mean((y_test - test_predictions_ridge)^2)
cat("Training Error (Ridge):", mse_train_ridge, "\n")

## Training Error (Ridge): 1.377282

cat("Test Error (Ridge):", mse_test_ridge, "\n")

## Test Error (Ridge): 1.554029

# Ridge regression (bootstrap)
bootstrap_mse <- function(data, indices, lambda) {
  # Create bootstrap sample
  bootstrap_sample <- data[indices, ]
  X_bootstrap <- as.matrix(bootstrap_sample[, -ncol(bootstrap_sample)])
  y_bootstrap <- bootstrap_sample$LC50

  model <- glmnet(X_bootstrap, y_bootstrap, alpha = 0, lambda = lambda)
  y_pred <- predict(model, s = lambda, newx = X_test)
  mse <- mean((y_test - y_pred)^2)
  return(mse)
}

# Bootstrap for multiple lambda values
bootstrap_results <- sapply(lambda_grid, function(lambda) {
  mse_values <- replicate(100, boot(trainData, bootstrap_mse, R = 1,
                                   lambda = lambda)$t)

  return(mean(mse_values))
})

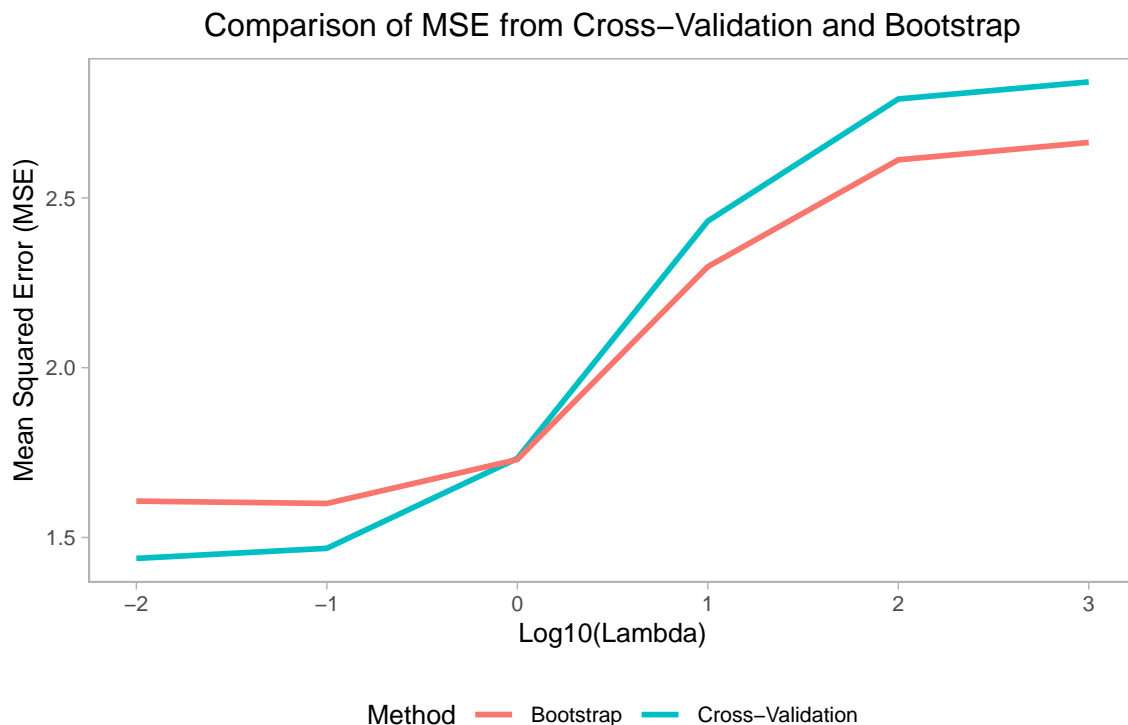
# Optimal Lambda from Bootstrap
optimal_lambda_bootstrap <- lambda_grid[which.min(bootstrap_results)]
cat("Optimal lambda from Bootstrap:", optimal_lambda_bootstrap, "\n")

## Optimal lambda from Bootstrap: 0.1

results_df <- data.frame(
  Lambda = lambda_grid,
  CV_MSE = sapply(lambda_grid,
                  function(l) mean(cv_ridge$cvm[cv_ridge$lambda == l])),
  Bootstrap_MSE = bootstrap_results
)

```

```
# Plot
ggplot(results_df, aes(x = log10(Lambda)) ) +
  geom_line(aes(y = CV_MSE, color = "Cross-Validation"), size = 1) +
  geom_line(aes(y = Bootstrap_MSE, color = "Bootstrap"), size = 1) +
  labs(title = "Comparison of MSE from Cross-Validation and Bootstrap",
       x = "Log10(Lambda)",
       y = "Mean Squared Error (MSE)",
       color = "Method") +
  theme_light() +
  theme(legend.position = "bottom",
        panel.grid = element_blank(),
        plot.title = element_text(hjust = 0.5),
        text = element_text(size = 10))
```



Non-linear Modeling Using Generalized Additive Models

```
sapply(trainData, function(x) length(unique(x)))
```

```
##   TPSA   SAacc   H050   MLOGP   RDCHI   GATS1p   nN   C040   LC50
##   171    160     9    283    259    289     9     6    347
```

```
library(mgcv)
```

```
set.seed(2024)
index <- sample(seq_len(nrow(data)), size = 2/3 * nrow(data))
trainData <- data[index, ]
testData <- data[-index, ]

X_train <- as.matrix(trainData[, -ncol(trainData)])
```

```

y_train <- trainData$LC50
X_test <- as.matrix(testData[, -ncol(testData)])
y_test <- testData$LC50

# GAM less complexity (k = 1)
k = 1
gam_model_1 <- gam(LC50 ~ s(TPSA, k=k) + s(SAacc, k=k) + s(H050, k=k) +
  s(MLOGP, k=k) + s(RDCHI, k=k) + s(GATS1p, k=k) +
  s(nN, k=k) + s(C040, k=k), data = trainData)

pred_gam_train_1 <- predict(gam_model_1, newdata = trainData)
pred_gam_test_1 <- predict(gam_model_1, newdata = testData)
mse_train_gam_1 <- mean((y_train - pred_gam_train_1)^2)
mse_test_gam_1 <- mean((y_test - pred_gam_test_1)^2)

cat("Training Error (GAM - k=1):", "\t", mse_train_gam_1, "\n")

## Training Error (GAM - k=1): 1.300005

cat("Test Error (GAM - k=1):", "\t", mse_test_gam_1, "\n")

## Test Error (GAM - k=1): 1.565394

# GAM more complexity (k = 6)
k = 6
gam_model_2 <- gam(LC50 ~ s(TPSA, k=k) + s(SAacc, k=k) + s(H050, k=k) +
  s(MLOGP, k=k) + s(RDCHI, k=k) + s(GATS1p, k=k) +
  s(nN, k=k) + s(C040, k=k), data = trainData)

pred_gam_train_2 <- predict(gam_model_2, newdata = trainData)
pred_gam_test_2 <- predict(gam_model_2, newdata = testData)
mse_train_gam_2 <- mean((y_train - pred_gam_train_2)^2)
mse_test_gam_2 <- mean((y_test - pred_gam_test_2)^2)

cat("Training Error (GAM - k=6):", "\t", mse_train_gam_2, "\n")

## Training Error (GAM - k=6): 1.147883

cat("Test Error (GAM - k=6):", "\t", mse_test_gam_2, "\n")

## Test Error (GAM - k=6): 1.634775

```

Regression Tree Model with Cost-Complexity Pruning

```

library(rpart)
library(rpart.plot)

tree_model <- rpart(LC50 ~ ., data = trainData, method = "anova")

train_predictions_tree1 <- predict(tree_model, newdata = trainData)
test_predictions_tree1 <- predict(tree_model, newdata = testData)

mse_train_tree1 <- mean((trainData$LC50 - train_predictions_tree1)^2)
mse_test_tree1 <- mean((testData$LC50 - test_predictions_tree1)^2)

```

```

printcp(tree_model)

##
## Regression tree:
## rpart(formula = LC50 ~ ., data = trainData, method = "anova")
##
## Variables actually used in tree construction:
## [1] GATS1p H050  MLOGP  nN      RDCHI  SAacc  TPSA
##
## Root node error: 949.5/364 = 2.6085
##
## n= 364
##
##      CP nsplit rel error  xerror    xstd
## 1  0.202576      0  1.00000 1.00521 0.081444
## 2  0.120003      1  0.79742 0.89217 0.074688
## 3  0.055504      2  0.67742 0.71621 0.060411
## 4  0.029750      4  0.56641 0.68528 0.057679
## 5  0.024807      5  0.53666 0.69253 0.060911
## 6  0.019644      6  0.51186 0.69168 0.061473
## 7  0.018758      7  0.49221 0.68003 0.061525
## 8  0.018127      8  0.47345 0.68003 0.061525
## 9  0.012598      9  0.45533 0.68531 0.063366
## 10 0.011999     10  0.44273 0.70330 0.069257
## 11 0.011431     11  0.43073 0.70041 0.069312
## 12 0.011033     13  0.40787 0.69407 0.069352
## 13 0.010592     14  0.39684 0.69141 0.068352
## 14 0.010000     15  0.38625 0.68556 0.068029

cat("Training Error (Tree):", mse_train_tree1, "\n")

## Training Error (Tree): 1.007531

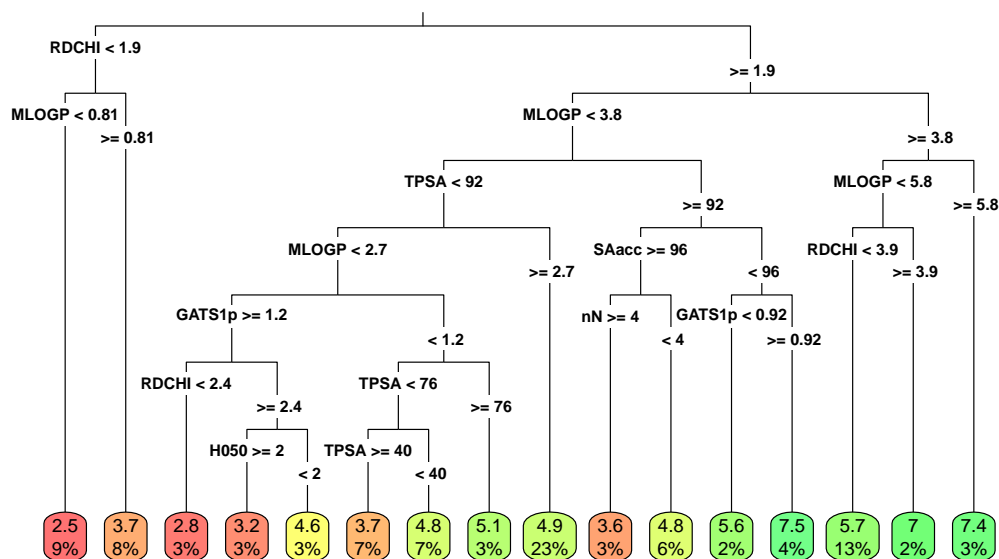
cat("Test Error (Tree):", mse_test_tree1, "\n")

## Test Error (Tree): 2.053224

rpart.plot(tree_model,
  main = "Customized Regression Tree for LC50",
  type = 3,
  fallen.leaves = TRUE,
  box.palette = "RdYlGn",
  cex = 0.6,
  tweak = 1,
  split.cex = 0.9,
  split.lty = 3,
)

```

Customized Regression Tree for LC50



```

optimal_cp <- tree_model$cptable[which.min(tree_model$cptable[, "xerror"]), "CP"]

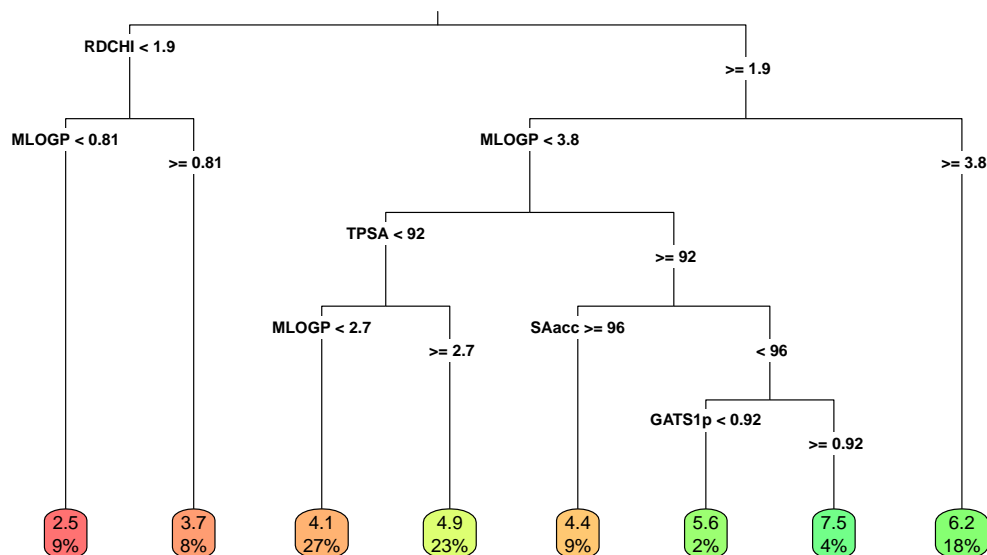
pruned_tree <- prune(tree_model, cp = optimal_cp)
train_predictions_tree <- predict(pruned_tree, newdata = trainData)
test_predictions_tree <- predict(pruned_tree, newdata = testData)

mse_train_tree <- mean((trainData$LC50 - train_predictions_tree)^2)
mse_test_tree <- mean((testData$LC50 - test_predictions_tree)^2)

rpart.plot(pruned_tree,
  main = "Pruned Regression Tree for LC50",
  type = 3,
  fallen.leaves = TRUE,
  box.palette = "RdYlGn",
  cex = 0.6,
  tweak = 1,
  split.cex = 0.9,
  split.lty = 3,
)

```

Pruned Regression Tree for LC50



```
printcp(pruned_tree)
```

```
##
## Regression tree:
## rpart(formula = LC50 ~ ., data = trainData, method = "anova")
##
## Variables actually used in tree construction:
## [1] GATS1p MLOGP RDCHI SAacc TPSA
##
## Root node error: 949.5/364 = 2.6085
##
## n= 364
##
##      CP nsplit rel error  xerror   xstd
## 1 0.202576      0  1.00000 1.00521 0.081444
## 2 0.120003      1  0.79742 0.89217 0.074688
## 3 0.055504      2  0.67742 0.71621 0.060411
## 4 0.029750      4  0.56641 0.68528 0.057679
## 5 0.024807      5  0.53666 0.69253 0.060911
## 6 0.019644      6  0.51186 0.69168 0.061473
## 7 0.018758      7  0.49221 0.68003 0.061525
```

```
cat("Training Error (Tree):", mse_train_tree, "\n")
```

```
## Training Error (Tree): 1.283949
```

```
cat("Test Error (Tree):", mse_test_tree, "\n")
```

```
## Test Error (Tree): 2.016152
```

Comparative Analysis of Regression Models

Training and Test Error Assessment for Various Models

```
library(tidyr)

model_comparison <- data.frame(
  Model = c("Linear", "Linear (dummy)", "Ridge", "GAM", "Tree"),
  Training_Error = c(error_train, error_train_dummy,
                     mse_train_ridge, mse_train_gam_2, mse_train_tree),
  Test_Error = c(error_test, error_test_dummy,
                 mse_test_ridge, mse_test_gam_2, mse_test_tree)
)

model_comparison_long <- pivot_longer(model_comparison,
                                       cols = c("Training_Error", "Test_Error"),
                                       names_to = "Error_Type",
                                       values_to = "MSE")

# Plot
ggplot(model_comparison_long, aes(x = Model, y = MSE, fill = Error_Type)) +
  geom_bar(stat = "identity", position = "dodge", alpha = 0.7) +
  labs(title = "Comparison of Training and Test Errors",
       x = "Model",
       y = "Mean Squared Error (MSE)",
       fill = "Error Type") +
  theme_light() +
  theme(legend.position = "bottom",
        panel.grid = element_blank(),
        plot.title = element_text(hjust = 0.5),
        text = element_text(size = 10))
```




```
ggplot(model_comparison_long, aes(x = Model, y = MSE,
                                color = Error_Type, group = Error_Type)) +
  geom_line(size = 0.9) +
  labs(title = "Line Plot Comparison of Training and Test Errors",
       x = "Model",
       y = "Mean Squared Error (MSE)",
       color = "Error Type") +
  theme_light() +
  theme(legend.position = "bottom",
        panel.grid = element_blank(),
        plot.title = element_text(hjust = 0.5),
        text = element_text(size = 10))
```



Classification Model Comparison for Diabetes Prediction

k-NN Classification and Cross-Validation Comparison

```
library(mlbench)
data(PimaIndiansDiabetes2)
df <- PimaIndiansDiabetes2
head(df)
```

```
##   pregnant glucose pressure triceps insulin mass pedigree age diabetes
## 1         6     148      72      35      NA 33.6   0.627  50      pos
## 2         1      85      66      29      NA 26.6   0.351  31      neg
## 3         8     183      64      NA      NA 23.3   0.672  32      pos
## 4         1      89      66      23     94 28.1   0.167  21      neg
## 5         0     137      40      35    168 43.1   2.288  33      pos
## 6         5     116      74      NA      NA 25.6   0.201  30      neg
```

```
na_count <- sapply(df, function(x) sum(is.na(x)))
print(na_count)
```

```
## pregnant glucose pressure triceps insulin mass pedigree age
##         0         5        35       227      374       11         0         0
## diabetes
##         0
```

KNN imputation

```
library(mlbench)
library(VIM)
```

```

df_imp <- kNN(df, k = 5)
df_imp <- df_imp[, 1:9]
df_imp$diabetes <- as.factor(df_imp$diabetes)
head(df_imp)

##   pregnant glucose pressure triceps insulin mass pedigree age diabetes
## 1         6      148      72      35      175 33.6    0.627  50      pos
## 2         1       85      66      29       55 26.6    0.351  31      neg
## 3         8      183      64      28      325 23.3    0.672  32      pos
## 4         1       89      66      23       94 28.1    0.167  21      neg
## 5         0      137      40      35      168 43.1    2.288  33      pos
## 6         5      116      74      27      112 25.6    0.201  30      neg

na_count_df_imp <- sapply(df_imp, function(x) sum(is.na(x)))
print(na_count_df_imp)

## pregnant glucose pressure triceps insulin mass pedigree age
##         0         0         0         0         0         0         0
## diabetes
##         0

library(caret)
set.seed(2024)
trainIndex <- createDataPartition(df_imp$diabetes,
                                   p = 2/3,
                                   list = FALSE,
                                   times = 1)

trainData <- df_imp[trainIndex, ]
testData <- df_imp[-trainIndex, ]

trainData$diabetes <- ifelse(trainData$diabetes == "pos", 1, 0)
testData$diabetes <- ifelse(testData$diabetes == "pos", 1, 0)
trainData$diabetes <- as.factor(trainData$diabetes)
testData$diabetes <- as.factor(testData$diabetes)

head(trainData)

##   pregnant glucose pressure triceps insulin mass pedigree age diabetes
## 1         6      148      72      35      175 33.6    0.627  50         1
## 2         1       85      66      29       55 26.6    0.351  31         0
## 3         8      183      64      28      325 23.3    0.672  32         1
## 5         0      137      40      35      168 43.1    2.288  33         1
## 9         2      197      70      45      543 30.5    0.158  53         1
## 12        10      168      74      32      171 38.0    0.537  34         1

library(class)
library(mlbench)
# Function: cross-validated k-NN errors
cv_knn_errors <- function(data, k_values, cv_type) {
  errors <- numeric(length(k_values))
  for (i in seq_along(k_values)) {
    k <- k_values[i]
    if (cv_type == "LOOCV") {
      control <- trainControl(method = "LOOCV")

```

```

    } else if (cv_type == "10-fold") {
      control <- trainControl(method = "cv", number = 10)
    }

    model <- train(diabetes ~ ., data = data, method = "knn",
                  trControl = control, tuneGrid = data.frame(k = k))
    errors[i] <- min(model$results$Accuracy)
  }
  return(1 - errors)
}

# k values (tested that >25 is too much)
k_values <- 1:25

# 5 fold
errors_10fold <- cv_knn_errors(df_imp, k_values, "10-fold")

# LOOCV
errors_loocv <- cv_knn_errors(df_imp, k_values, "LOOCV")

trainIndex <- createDataPartition(df_imp$diabetes, p = .67, list = FALSE)
trainData <- df_imp[trainIndex, ]
testData <- df_imp[-trainIndex, ]

# Loop for each k
test_errors <- numeric(length(k_values))
for (i in seq_along(k_values)) {
  k <- k_values[i]
  predictions <- knn(train = trainData[, -ncol(trainData)],
                     test = testData[, -ncol(testData)],
                     cl = trainData$diabetes, k = k)
  test_errors[i] <- mean(predictions != testData$diabetes)
}

# Create a data frame for plotting
error_data <- data.frame(
  k = k_values,
  LOOCV_Error = errors_loocv,
  Fold10_Error = errors_10fold,
  Test_Error = test_errors
)

error_data_long <- reshape2::melt(error_data, id.vars = "k")

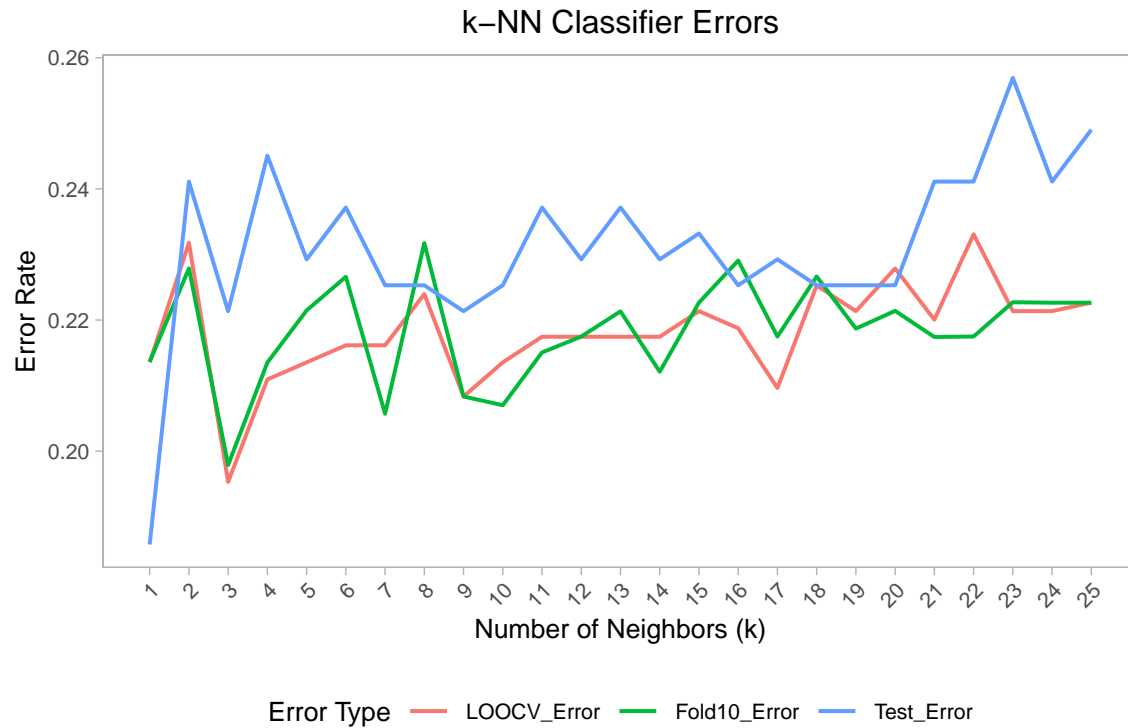
# Plotting
ggplot(error_data_long, aes(x = k, y = value, color = variable)) +
  geom_line(size = 0.7) +
  # geom_point(size = 1.3) +
  labs(title = "k-NN Classifier Errors",
       x = "Number of Neighbors (k)",
       y = "Error Rate",
       color = "Error Type") +
  scale_x_continuous(breaks = k_values) +
  theme_light() +

```

```

theme(axis.text.x = element_text(angle = 45, hjust = 1),
      legend.position = "bottom",
      panel.grid = element_blank(),
      plot.title = element_text(hjust = 0.5),
      text = element_text(size = 10))

```



Generalized Additive Model for Diabetes Prediction

```

library(caret)
library(mgcv)
library(dplyr)
library(broom)

# model
gam_model <- gam(diabetes ~ s(pregnant) + s(glucose) + s(pressure) +
                 s(triceps) + s(insulin) + s(mass) +
                 s(pedigree) + s(age), data = trainData, family = binomial)
summary(gam_model)

```

```

##
## Family: binomial
## Link function: logit
##
## Formula:
## diabetes ~ s(pregnant) + s(glucose) + s(pressure) + s(triceps) +
##           s(insulin) + s(mass) + s(pedigree) + s(age)
##
## Parametric coefficients:

```

```
##           Estimate Std. Error z value Pr(>|z|)
## (Intercept) -1.3043      0.1805  -7.225   5e-13 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approximate significance of smooth terms:
##           edf Ref.df Chi.sq  p-value
## s(pregnant) 1.000  1.001  8.278 0.004022 **
## s(glucose)  4.196  5.166 21.735 0.000627 ***
## s(pressure) 1.000  1.000  1.014 0.314007
## s(triceps)  1.107  1.206  2.598 0.112210
## s(insulin)  8.411  8.889 33.158 8.95e-05 ***
## s(mass)      3.916  4.866 16.355 0.005263 **
## s(pedigree) 1.000  1.000  0.001 0.971580
## s(age)       2.222  2.806  5.477 0.131900
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) =  0.435   Deviance explained = 41.3%
## UBRE = -0.14736   Scale est. = 1           n = 515

# library(MASS)
# # full model
# full_model <- gam(diabetes ~ s(pregnant) + s(glucose) + s(pressure) +
#                   s(triceps) + s(insulin) + s(mass) +
#                   s(pedigree) + s(age),
#                   data = trainData, family = binomial)
# # stepwise AIC
# selected_model <- step(full_model, direction = "backward")
# summary(selected_model)
```

Tree-Based Methods for Diabetes Classification

```
library(rpart)
library(ipred)
library(randomForest)
library(caret)

classification_tree <- rpart(diabetes ~ ., data = trainData, method = "class")

tree_train_pred <- predict(classification_tree, trainData, type = "class")
tree_test_pred <- predict(classification_tree, testData, type = "class")

# error
tree_train_error <- mean(tree_train_pred != trainData$diabetes)
tree_test_error <- mean(tree_test_pred != testData$diabetes)

cat("Classification Tree Training Error:", tree_train_error, "\n")

## Classification Tree Training Error: 0.1359223
cat("Classification Tree Test Error:", tree_test_error, "\n")

## Classification Tree Test Error: 0.201581
```

```

bagged_trees <- bagging(diabetes ~ ., data = trainData)

bagged_train_pred <- predict(bagged_trees, trainData)
bagged_test_pred <- predict(bagged_trees, testData)

# error
bagged_train_error <- mean(bagged_train_pred != trainData$diabetes)
bagged_test_error <- mean(bagged_test_pred != testData$diabetes)

cat("Bagged Trees Training Error:", bagged_train_error, "\n")

## Bagged Trees Training Error: 0.001941748

cat("Bagged Trees Test Error:", bagged_test_error, "\n")

## Bagged Trees Test Error: 0.1936759

# Fit the random forest
random_forest <- randomForest(diabetes ~ ., data = trainData)

# Predict on training and test data
rf_train_pred <- predict(random_forest, trainData)
rf_test_pred <- predict(random_forest, testData)

# Calculate training and test error
rf_train_error <- mean(rf_train_pred != trainData$diabetes)
rf_test_error <- mean(rf_test_pred != testData$diabetes)

cat("Random Forest Training Error:", rf_train_error, "\n")

## Random Forest Training Error: 0

cat("Random Forest Test Error:", rf_test_error, "\n")

## Random Forest Test Error: 0.1818182

# Store the errors in a data frame
results <- data.frame(
  Model = c("Classification Tree", "Bagged Trees", "Random Forest"),
  Training_Error = c(tree_train_error, bagged_train_error, rf_train_error),
  Test_Error = c(tree_test_error, bagged_test_error, rf_test_error)
)

# Reshape the data for ggplot
results_long <- reshape2::melt(results, id.vars = "Model", variable.name = "Error_Type", value.name = "Error_Rate")

ggplot(results_long, aes(x = Model, y = Error_Rate, fill = Error_Type)) +
  geom_bar(stat = "identity", position = position_dodge(width = 0.9),
    width = 0.7) +
  geom_text(aes(label = round(Error_Rate, 3)),
    position = position_dodge(width = 0.9),
    vjust = -0.5, size = 2) + # Data labels
  theme_light() +
  labs(title = "Comparison of Training and Test Errors",
    x = "Model",
    y = "Error Rate",
  )

```

```

fill = "Error Type") +
theme(legend.position = "bottom",
      panel.grid = element_blank(),
      plot.title = element_text(hjust = 0.5),
      axis.title = element_text(),
      text = element_text(size = 10))

```



Neural Network Modeling for Diabetes Prediction

```

library(nnet)
library(caret)

neurons <- seq(1, 100, by = 5)

# Initializing
train_errors <- c()
test_errors <- c()

# Loop
for (n in neurons) {
  # Fit NN
  nn_model <- nnet(diabetes ~ ., data = trainData,
                   size = n, linout = FALSE, maxit = 500, trace = FALSE)

  # Train
  train_pred <- predict(nn_model, trainData)
  train_pred_class <- ifelse(train_pred > 0.5, 1, 0) # Convert to binary
  train_error <- mean(train_pred_class != trainData$diabetes)
}

```



```

train_errors <- c(train_errors, train_error)

# Test
test_pred <- predict(nn_model, testData)
test_pred_class <- ifelse(test_pred > 0.5, 1, 0) # Convert to binary
test_error <- mean(test_pred_class != testData$diabetes)
test_errors <- c(test_errors, test_error)
}

# Results
error_results <- data.frame(
  Neurons = neurons,
  Training_Error = train_errors,
  Test_Error = test_errors
)
print(error_results)

```

```

##      Neurons Training_Error Test_Error
## 1         1             1           1
## 2         6             1           1
## 3        11             1           1
## 4        16             1           1
## 5        21             1           1
## 6        26             1           1
## 7        31             1           1
## 8        36             1           1
## 9        41             1           1
## 10       46             1           1
## 11       51             1           1
## 12       56             1           1
## 13       61             1           1
## 14       66             1           1
## 15       71             1           1
## 16       76             1           1
## 17       81             1           1
## 18       86             1           1
## 19       91             1           1
## 20       96             1           1

```

```

# Choose the optimal number of neurons based on test error
optimal_neurons <- neurons[which.min(test_errors)]
cat("Optimal number of neurons:", optimal_neurons, "\n")

```

```
## Optimal number of neurons: 1
```

```

library(ggplot2)
error_results_long <- reshape2::melt(error_results,
                                     id.vars = "Neurons",
                                     variable.name = "Error_Type",
                                     value.name = "Error")

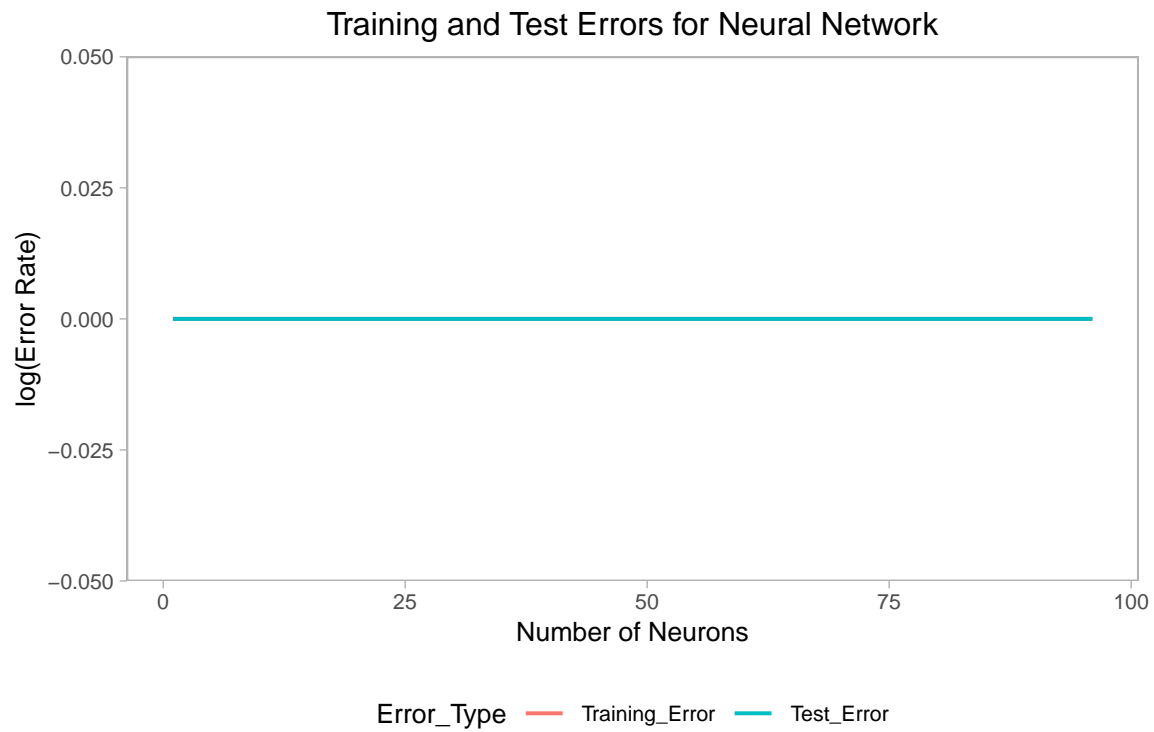
ggplot(error_results_long, aes(x = Neurons, y = log10(Error),
                              color = Error_Type)) +
  geom_line(size = 0.7) +
  labs(title = "Training and Test Errors for Neural Network",

```

```

x = "Number of Neurons",
y = "log(Error Rate)" +
theme_light() +
theme(legend.position = "bottom",
      panel.grid = element_blank(),
      plot.title = element_text(hjust = 0.5),
      text = element_text(size = 10))

```



Comparing Classification Methods for Diabetes Analysis